

PRESEARCH.md

Date: 2026-02-16

Source: `mvp-1-collab-board/G4 Week 1 - CollabBoard-requirements.pdf`

1) Problem and Constraints

We need to ship a real-time collaborative whiteboard with an AI board agent in one sprint.

Hard deadlines:

- Pre-Search checkpoint: Monday, 2026-02-16 (first hour)
- MVP checkpoint: Tuesday, 2026-02-17 (24 hours)
- Early submission target: Friday, 2026-02-20
- Final deadline: Sunday, 2026-02-22, 10:59 PM CT

MVP hard gate (must all pass):

- Infinite board with pan/zoom
- Sticky notes with editable text
- At least one shape type
- Create/move/edit objects
- Real-time sync (2+ users)
- Multiplayer cursors with labels
- Presence awareness
- User authentication
- Public deployment

2) Phase 1: Define Constraints

2.1 Scale and Load Profile (assumptions for sprint)

- Launch: 5-20 concurrent users per board.
- 6 months: 100-500 weekly active users if project extends.
- Traffic pattern: spiky (class demos and review sessions).
- Real-time requirements: object sync target <100ms, cursor sync target <50ms.
- Cold start tolerance: low for collaboration path, medium for AI command path.

2.2 Budget and Cost Ceiling

- Sprint dev/testing budget target: <= \$150.
- Early production target (first 1,000 users): <= \$700 per month.
- Tradeoff: pay managed infra costs to reduce implementation risk.

2.3 AI Cost Modeling (required)

Assumptions:

- Commands per session: 6
- Sessions per user per month: 8
- Command mix: 80% simple, 20% complex
- Token model:
 - simple command: 900 input + 300 output
 - complex command: 1600 input + 1200 output
- Weighted average tokens per command: 1520
- Tokens per user per month: $6 * 8 * 1520 = 72,960$
- Blended LLM cost assumption: \$3.20 per 1M tokens (planning value, provider-adjustable)

Projected monthly costs:

Scale	LLM Tokens / Month	LLM Cost	Infra Cost (hosting/db/functions)	Total
100 users	7.296M	\$23	\$90	\$113
1,000 users	72.96M	\$233	\$220	\$453
10,000 users	729.6M	\$2,335	\$1,250	\$3,585
100,000 users	7.296B	\$23,347	\$7,500	\$30,847

2.4 Time to Ship

- Primary priority this week: speed-to-market with stable multiplayer.
- Maintainability guardrails: typed contracts, decision log, automated tests.
- Iteration cadence: daily cut and checkpoint review.

2.5 Compliance and Regulatory

- No healthcare scope in MVP.
- Baseline privacy: least-privilege rules and no secrets in client runtime.
- Accessibility baseline (sprint scope): keyboard operability, visible focus, and contrast evidence.
- Note: this accessibility evidence is also relevant for regulated/public-sector evaluations.
- If later targeting enterprise/government deployment: SOC 2 controls, ATO workflows, and formal audit packages.

2.6 Team and Skills

- Execution baseline: TypeScript-first.
- Best velocity stack this week: React + Firebase + server-side AI command orchestration.

3) Phase 2: Architecture Discovery

Option Comparison

Option	Stack	Pros	Cons	Fit for 1-week sprint
A	React + Konva + Firebase (Auth + Firestore + RTDB presence) + Cloud Functions	Fastest setup, managed auth, realtime primitives	Vendor lock-in, Firestore modeling discipline needed	Best
B	React + Konva + Supabase (Auth + Postgres + Realtime) + Edge Functions	SQL flexibility, good DX	More realtime conflict plumbing for board semantics	Good
C	React + custom WebSocket + Redis + Postgres	Maximum control and tuning	Highest implementation and ops risk	Poor

Selected option: A.

3.0 Build-vs-Buy Scorecard (OSS-First)

Goal: avoid paid lock-in unless speed/risk clearly requires it.

Option Comparison

Option	License/Cost	Strengths	Risks	Recommendation
Current: Firebase + custom sync	Managed usage-based, generous free tier	Already implemented, fastest continued delivery	Not CRDT by default, vendor coupling	Keep for MVP + immediate iteration
Yjs + Hocuspocus + Postgres/Supabase	Open source core (self-hostable)	Strong CRDT collaboration model, no per-seat vendor tax	Higher ops complexity and migration effort	Best OSS migration candidate post-MVP
Automerger + custom transport	Open source core	Full control, strong local-first story	Highest engineering effort for production hardening	R&D only unless team size increases
Liveblocks	Commercial SaaS (free starter tier)	Fastest advanced collaboration features (comments, presence, threads)	Ongoing SaaS spend and dependency	Use only if deadline risk spikes

Decision checkpoint:

- Current phase (through MVP + submission): stay on Firebase custom sync.
- Post-MVP phase (starting 2026-02-23): run a 1-2 day Yjs/Hocuspocus spike and compare migration cost against roadmap value.
- Escalate to Liveblocks only if we miss reliability targets under >20 concurrent users or cannot ship required collaboration features on schedule.
- Spike output artifact: `YJS_SPIKE.md`.

3.1 Hosting and Deployment

- Frontend: Firebase Hosting.
- API and AI actions: Firebase Cloud Functions.
- CI: GitHub Actions for lint, tests, and deploy checks.

3.2 Authentication and Authorization

- Primary MVP auth: Firebase Auth with Google OAuth.
- Fallback if blocked: Firebase email-link.
- Board access control enforced in Firestore rules.

3.3 Data Model Contracts (explicit)

Canonical object schema:

```
```ts
type BoardObject = {
 id: string
 boardId: string
 type: 'stickyNote' | 'shape' | 'text' | 'frame' | 'connector'
 position: { x: number; y: number }
 size?: { width: number; height: number }
 rotation?: number
 zIndex: number
 color?: string
 text?: string
 shapeType?: 'rectangle' | 'circle' | 'line'
 fromId?: string
 toId?: string
 createdBy: string
 createdAt: number
 updatedBy: string
 updatedAt: number
 version: number
 deleted?: boolean
}..
```

Presence and cursor schema:

```
```ts
type CursorPresence = {
  boardId: string
  userId: string
  displayName: string
  color: string
  x: number
  y: number
  lastSeen: number
  connectionId: string
}..
```

Storage layout:

- Firestore:
 - `boards/{boardId}` metadata
 - `boards/{boardId}/objects/{objectId}` canonical objects
 - `boards/{boardId}/aiCommands/{commandId}` idempotency and status records
- Realtime Database:
 - `presence/{boardId}/{userId}` ephemeral online status and cursor coordinates

3.4 Conflict Resolution Strategy (explicit)

- Conflict model: Last-Write-Wins (LWW) using authoritative server `updatedAt`.
- Every object write increments `version` and sets `updatedAt`.
- Client UX uses optimistic updates, then reconciles to server state if overwritten.

ten.

- Text edits commit on blur/submit (not every keystroke) to reduce collision rate.
- Multi-object AI writes use batched writes; dependent edits use transaction checks where needed.

3.5 Sync Performance Model (explicit)

- Cursor sync: throttle publishes to <=20 updates/sec/user (~50ms interval target).
- Object drag sync: throttle to <=10 updates/sec/object while dragging, plus final commit on pointer up.
- Presence heartbeat: update `lastSeen` every 10-15s.
- Listener topology per board:
 - one object listener for `boards/{boardId}/objects`
 - one presence listener for `presence/{boardId}`
- Index strategy:
 - rely on single-field indexes for MVP subcollection access
 - add composite index only if sort/filter combinations require it
- Batched writes for template generation keep write bursts bounded and predictable.

3.6 AI Agent Execution Model (explicit)

Request path:

1. Client sends command with `clientCommandId`.
2. Cloud Function validates command and checks idempotency store.
3. Server fetches board context via `getBoardState()` (MVP cap: latest 500 objects summary).
4. LLM planner returns tool-call plan.
5. Server executes tool calls sequentially and writes updates.
6. Command result persisted in `aiCommands` and shared to all users through normal board sync.

Concurrency and safety:

- Dispatcher exposes rubric tool functions: `createStickyNote`, `createShape`, `createFrame`, `createConnector`, `moveObject`, `resizeObject`, `updateText`, `changeColor`, `getBoardState`.
- Tool calls execute server-side only (no direct client key usage).
- Idempotency key: `clientCommandId` per board.
- Simultaneous AI commands: FIFO per-board queue semantics.
- If two users submit commands concurrently, ordering is deterministic by command creation timestamp.

3.7 Offline and Reconnect Strategy

- Enable Firestore web offline persistence.
- Use RTDB `onDisconnect()` to remove stale presence entries.
- On reconnect:
 - display `Reconnecting` and `Syncing` UI states
 - replay local pending writes via Firestore client
 - refresh board state from canonical object listener

3.8 Frontend Rendering Risk and Mitigation

Risk:

- Naive React + Konva re-renders can miss 60 FPS at 500+ objects.

Mitigation:

- Keep hot interaction state in an imperative stage manager (refs/store), not full React tree updates.
- Use React state for controls and metadata, not per-frame pointer movement.

3.9 Testing Tooling

- Unit: Vitest for reducers/transforms/command parsing.
- Integration: Firebase Emulator Suite for auth/rules/sync flows.

- E2E: Playwright with two browser contexts and scripted multi-user scenarios.

4) Phase 3: Post-Stack Refinement

Security Risks and Mitigations

- Risk: over-permissive database rules.
 - Mitigation: deny-by-default rules and explicit board membership checks.
- Risk: prompt injection via AI command text.
 - Mitigation: strict tool schema validation and no arbitrary code execution.
- Risk: leaked API keys.
 - Mitigation: server-only keys in environment secrets.

Accessibility Artifacts

- `docs/ACCESSIBILITY_AUDIT.md`: checklist and evidence log for keyboard/focus/contrast/ARIA.
- `docs/VPAT_DRAFT.md`: optional VPAT-style conformance report scaffold.

Project Structure

- `apps/web/` React whiteboard app
- `apps/functions/` AI and server workflows
- `packages/shared/` shared types and schemas

Naming and Style

- TypeScript strict mode.
- ESLint + Prettier.
- Conventions: `camelCase` vars/functions, `PascalCase` components/types.

Post-MVP Feature Roadmap (Prioritized)

AI-first differentiators (highest leverage):

- Smart layout analysis and auto-arrange commands.
- Sticky/theme synthesis from board content.
- Layout suggestions based on board intent (retro/roadmap/brainstorm).
- Multi-modal ingest (screenshot to stickies via OCR).
- One-shot template generation with multi-step AI plans.

Interactive parity features:

- Connectors/arrows with snapping.
- Frames and grouped regions with titles.
- Undo/redo command history.
- Comments with mentions.
- Voting mode and facilitation timer.
- Mini-map navigation.

Unique extensions:

- Voice commands (Web Speech API).
- Template library with one-click board setups.
- Export selection/full board to image and PDF.
- Activity timeline (board replay).
- AI chat sidebar grounded in current board state.

Polish:

- Keyboard shortcut reference panel.
- First-run onboarding tour.
- Loading skeleton states for async operations.
- Lightweight object animation for create/delete/move feedback.

5) Final Stack Decision

- Frontend: React + TypeScript + Konva
- Realtime and auth: Firebase (Firestore + RTDB + Auth)
- AI: function-calling through server-side dispatcher
- Hosting: Firebase Hosting + Cloud Functions

Why this stack now:

- Minimum integration overhead for 24-hour MVP gate.
- Fastest path to stable collaboration and auth.
- Lock-in risk accepted for short-term delivery confidence.